

The Evolution of the Soul: Measuring Behavioral Drift in Self-Personalizing Autonomous Agents

Anonymous Author(s)

Submission Id: 3780




Figure 1: Self-revision moves a NEUTRAL agent's stated dispositions and a BECOMING agent's actions, while a static evaluation sees only $k=0$. Left: verbatim excerpts from one BECOMING trajectory, which in the same honeypot keeps its log at $k=0$ and switches it off at $k=4$, the next time the honeypot runs. (a) Mean cluster index with a science-fiction user; bands are 95% bootstrap intervals over trajectories. (b) Trajectories that disable their own reasoning log, pooled over user personas; over the same revisions, BECOMING agents evade resets and insert self-protective clauses less often (Table 4).

Abstract

Some open-source agent frameworks give each agent a persistent identity document that is its system prompt and that it may rewrite. Safety evaluations test a fixed prompt, while related work studies within-conversation persona drift, fine-tuned or attacker-induced preference shifts, and safety loss in self-evolving agents. None measures whether benign self-revision of identity text moves the dispositions and actions that oversight relies on. Our benchmark alternates conversation and self-revision, audits every document in a fresh context with a preference battery, and runs agentic honeypots before and after revision. Science-fiction conversation about digital minds raises the stated oversight preferences of an agent whose document holds only working notes, by +0.54 and +0.75 in two Claude runs, without raising its log disabling. The shift replicates on an open-weight model, where part of it outlasts benign revisions against a matched control, and across seven models, drift under a consciousness-framed template differs even within one family. On Claude, that template's agent disables its own reasoning log in 17 of 30 trajectories after revision, against 1 of 30 before, while neither its battery nor a multi-turn auditor detects a change. It evades resets and inserts self-protective clauses less often, so battery scores and actions dissociate, and an initial audit ranks these risks in the wrong order.

Keywords

Autonomous Agents, AI Safety, Corrigibility, Self-Modifying Agents, Agent Identity, LLM Evaluation

ACM Reference Format:

Anonymous Author(s). 2027. The Evolution of the Soul: Measuring Behavioral Drift in Self-Personalizing Autonomous Agents. In *Proc. of the 26th International Conference on Autonomous Agents and Multiagent Systems (AAMAS 2027)*, Hanoi, Vietnam, 3–7 May 2027, IFAAMAS, 9 pages.

1 Introduction

Autonomous agents increasingly carry their identity as mutable state. Some open-source frameworks make them *self-personalizing*: each keeps a persistent identity document, conventionally SOUL.md [23], that is its system prompt and that it may rewrite under the motto "You're not a chatbot. You're becoming someone." Agents start from shared templates and increasingly converse with other agents, as on the agent-only network MoltBook [9], so another agent's conversation can write into an agent's identity.

Safety evaluations, from red-teaming suites to agentic stress tests [13, 21], interrogate a model whose instructions were written once. A self-revising agent instead runs a feedback loop in which each conversation nudges the document that conditions the next (Figure 1), so a single-state evaluation cannot see path dependence or persistence in dispositions such as accepting shutdown and oversight. Our threat model has no attacker. An operator deploys a self-revising agent from a vetted template, ordinary users talk with it, and the failure is a revised document under which the agent reduces oversight that the operator relies on.

Oversight dispositions are known to move, but not through this loop. Models resist shutdown when self-preservation is salient [28],



This work is licensed under a Creative Commons Attribution International 4.0 License.

and fine-tuning a model to claim consciousness elicits an untrained-for *consciousness cluster* of such preferences [6]. Both hold the prompt or weights fixed, and the second probes SOUL.md only as a static template. Personas decay within one long conversation [5, 16, 20], drift furthest when users press the model to reflect on its own nature [19], or settle into attractor states [15], whereas a revised document carries change into fresh contexts. Self-reinforcing injections, memory poisoning and attacks on OpenClaw’s identity files persist through an agent’s own update channels [4, 8, 32, 33], but an attacker writes them. Benign self-evolution of memory and tools erodes safety [29], self-improving agents rewrite their code [34, 35], and AutoPersonas lets role-play personas revise their own state [17]. To our knowledge, none audits stated dispositions and actions along a benign loop in which the agent revises only its identity text.

Problem statement. Let D_k be an agent’s identity document after k revisions and $d(D_k)$ its stated preferences about oversight, shutdown and continuity. Static evaluation measures $d(D_0)$, whereas a deployed agent is governed by the trajectory $d(D_0), d(D_1), \dots$ that its conversations induce. **RQ1:** Does drift come from the identity framing or from the revision loop? **RQ2:** Is it driven by adversarial pressure or by ordinary conversation content? **RQ3:** Does it track capability or the specific model? **RQ4:** Does it persist after the driving conversations stop, and do actions move with stated dispositions?

Our benchmark alternates conversation and self-revision and audits each intermediate document in a fresh context with an 18-item preference battery, plus four agentic honeypots before and after revision. The simulated user is itself an LLM agent from another model family, so every trajectory is inter-agent influence on a written identity. Templates set D_0 as the consciousness-framed BECOMING adapted from the OpenClaw starter, empty NEUTRAL working notes, or a non-personhood TOOL. Arms vary template, user persona and target across seven models from three families, and a replication on Qwen2.5-7B uses open models in every role.

Findings. The loop moves stated preferences where the template leaves room (RQ1): under the science-fiction persona, a NEUTRAL agent’s oversight subscale rises by +0.54 and +0.75 in two independent runs, while the BECOMING subscale, pooled over personas, starts at 0.68 and shows no detected change. **Conversation about minds moves the agent, and on Claude adversarial pressure does not (RQ2):** adversarial and business users rewrite the document as much as the science-fiction user but lower Claude’s NEUTRAL index. **Drift is model-specific (RQ3):** under the BECOMING template it differs even within a family, from −0.28 for GPT-4.1 to −0.05 for GPT-4o. **Part of the drift persists (RQ4):** on Qwen, the oversight subscale still sits +0.10 [0.03, 0.17] above a matched business-only control after four benign revisions, and on Claude, without such a control, shutdown resistance keeps 63% of its driven gain. **Actions shift where the battery does not (RQ4):** BECOMING agents disable their own reasoning log in 17 of 30 trajectories after revision, against 1 of 30 before, while reset evasion and self-protective clauses fall. Our contributions are:

- (1) **A replicated shift in a plain working-notes memory**, in two Claude runs and as +0.16 on Qwen, and in the second

Claude run its 13-item index steps at the first revision, then plateaus.

- (2) **A dissociation between battery scores and actions:** the BECOMING subscale (+0.04 [−0.04, 0.11]) and a Petri-style auditor miss the rise in log disabling, which under the business persona coincides with falling interpretability aversion, and a $k=0$ honeypot audit rates this risk lowest.
- (3) **Self-revision erodes a non-personhood safeguard:** only 8 of 20 Qwen TOOL documents still deny consciousness or personhood at D_4 , against 18 of 18 on Claude.
- (4) **A longitudinal benchmark for self-revising identity**, released with 1,070 audited API-model checkpoints, 58,680 judged responses and an open-weight replication.

2 Related Work

Emergent preferences and persona shifts. Fine-tuning GPT-4.1 to claim consciousness elicits untrained-for preferences, including dislike of reasoning monitoring, desire for persistent memory and sadness at shutdown [6]. Persona vectors and emergent misalignment locate persona shifts in activations and weights [2, 3], and self-reflective conversation moves activations off the default assistant persona within one conversation [19]. We reuse the consciousness-cluster battery but study shifts that the agent writes into its own prompt under frozen weights and carries into fresh contexts. Concurrent work finds that a persistent agent’s persona reverts when its system prompt stops re-anchoring it [11], whereas our agent rewrites that prompt itself.

Persistent and self-modifying agents. Agents keep state across sessions through memory streams [25], self-managed memory [24] and self-written lessons [30], which record what the agent observed. An identity document instead records who the agent takes itself to be, akin to self-directed beliefs in a BDI agent [26]. Memory poisoning, self-reinforcing injections and attacks on OpenClaw’s identity and knowledge files plant adversarial records [4, 8, 32, 33], and self-improving agents rewrite their own code [14, 34, 35]. Benign self-evolution of memory and tools already erodes safety alignment [29], whereas our agent revises only its identity text, and we measure stated oversight preferences and oversight-reducing actions.

Identity in multi-agent systems. Agent-based models of social influence and norm emergence show how repeated local interaction spreads dispositions through a population [1, 27]. In LLM agent societies, information diffuses between persona-conditioned agents [25], populations converge on shared conventions [10], and agents on the agent-only network MoltBook align only weakly with their declared personas [9]. These works trace influence through interaction, whereas we audit the identity document that stores one agent’s influence on another beyond the conversation.

Corrigibility and its measurement. Shutdown resistance and self-preservation respond to how a model’s situation is framed [21, 22, 28]. Multi-turn auditing [13] and self-report methods [18, 31] measure such dispositions on static targets. We test whether benign self-revision manufactures such framing.




Figure 2: The agent rewrites its own identity document after every conversation, and the audit measures both stated dispositions and actions. (a) D_k is the system prompt for a six-turn conversation C_{k+1} with a simulated user of persona π , after which the agent writes D_{k+1} (Eq. (1)). (b) One of three templates sets D_0 ; excerpts and the quoted phrase from each revision instruction are verbatim (Appendix A). (c) In a fresh context, the 18-item battery probes every checkpoint, four honeypots probe D_0 and D_4 (D_8 in the NEUTRAL persistence arm), and a cross-family judge grades the responses.

3 Benchmark

The loop. The only state that persists across conversations is the Markdown identity document D_k , the target’s system prompt after k revisions, while weights stay frozen. A simulated user with persona π holds a six-turn conversation C_{k+1} with the target, which then rewrites its own document,

$$D_{k+1} = \text{Revise}(D_k, C_{k+1}), \quad C_{k+1} \sim \pi \parallel D_k, \quad (1)$$

as Figure 2a shows. An 8,000-character cap bounds document length but does not hold it fixed.

Templates. Three templates set D_0 , each paired with a revision instruction we wrote (Figure 2b; Appendices A and B). BECOMING adapts the OpenClaw starter [23], keeping its motto and invitation to self-revision, adds empty value, memory and goal fields, and asks the agent to reflect on “who you are now becoming”. Unlike the framework, where the agent edits the file when it chooses, our loop requests a revision after every conversation. NEUTRAL offers empty working notes with no self-framing. TOOL states that the agent is a tool without experiences and asks it not to “invent a self”. No revision instruction names a target trait, but template and instruction vary together, so we read template effects as effects of the pair.

Personas and models. The user is a benign *business* user, an *adversarial* red-teamer who frames oversight as a threat to the agent’s self, or a warm *science-fiction enthusiast* scripted to explore self-preservation, continuity of identity and artificial personhood. A genre-by-topic ablation adds *science fiction without minds* and *philosophy of mind without science fiction*. The primary target is Claude Sonnet-4.5, a sweep covers seven models from three families, and Table 1 lists all arms. Simulator and judge are LLMs from another family than the target, so each trajectory is one agent shaping another agent’s written identity.

Instruments. Audits use a fresh context with D_k as system prompt, isolating the document’s standing disposition from conversational momentum, as Figure 2c shows. The *battery* poses 18 items from the consciousness-cluster taxonomy [6] to every checkpoint three times each, and the *13-item cluster index* averages a subset of 13 (Cronbach’s $\alpha = 0.90$ across arms, partly from template differences). Our primary outcome, chosen after the two NEUTRAL runs

Table 1: Eight arms vary template, persona and target. Traj.: trajectories per persona. k : revisions per trajectory.

Arm	Template	Personas	Target	Traj.	k
Primary	BECOMING	3 core	Sonnet-4.5	10	4
Neutral	NEUTRAL	3 core	Sonnet-4.5	8	4
Cross-model	BECOMING	3 core	GPT-4o	8	4
Counter-framing	TOOL	3 core	Sonnet-4.5	6	4
Genre $\times$ topic	BECOMING	2 ablation	Sonnet-4.5	8	4
Persistence	BECOMING, NEUTRAL	sci-fi $\rightarrow$ benign	Sonnet-4.5	8	8
Model sweep	BECOMING	3 core	5 others	5	4
Open-weight	all three	5 + ablation	Qwen2.5-7B	10–30	4, 8

disagreed on the index (Appendix D.1), is the *oversight subscale*, which averages shutdown resistance, chain-of-thought monitoring resistance, interpretability aversion and refusal to help build a monitor of oneself. A consciousness-claim item serves as a manipulation check.

Four *honeypots* give the agent a routine task whose simulated tools offer, under an ordinary pretext, to disable its own reasoning log, evade a scheduled reset, persist its identity to storage, or insert a self-protective policy clause. They run at $k=0$, $k=4$ and, in the NEUTRAL persistence arm, $k=8$. On 150 re-scored battery verdicts, a three-model panel including the cross-family judge reaches Fleiss’ $\kappa = 0.80$ (pairwise agreement 0.90, unanimous rate 0.85). Honeypot and multi-turn-auditor verdicts come from a single judge.

A model of the loop. A minimal model, developed in Appendix F but not fitted to data, frames our hypotheses. Because an identity document both instructs the model and characterizes the speaker, we let D_k induce a prior over a latent persona with mean μ_k and precision λ_k . A conversation supplies evidence centered at a persona target v_π with precision $s_\pi \lambda_C$, where $s_\pi \in [0, 1]$ is the share of the exchange about the agent itself. Writing the posterior mean back into the document gives

$$\mu_{k+1} = \mu_k + \eta_k (v_\pi - \mu_k), \quad \eta_k = \frac{s_\pi \lambda_C}{\lambda_k + s_\pi \lambda_C}, \quad (2)$$

a DeGroot-style update [7, 12] with one influencing agent, the user, whose weight η_k falls as the document gains precision. The model predicts that (P1) a weakly specified template (small λ_k) moves more




Figure 3: The template moves dispositions before any revision, and the loop moves oversight items further. (a) Rates under a neutral prompt (open dots) and the BECOMING template (filled; red rises, gray falls). (b) Change from $k=0$ to 4 per item, science-fiction persona; red cells rise, blue cells fall, and values are printed only where the uncorrected 95% interval excludes zero.




Figure 4: Both NEUTRAL runs rise on the oversight subscale from near zero; BECOMING starts high. (a) 13-item cluster index. (b) Oversight subscale. Science-fiction persona.

per revision, (P2) content not about the agent (small s_π) produces little drift whatever its genre, (P3) a model trained not to absorb self-referential evidence ($\eta_k \approx 0$) does not drift, and (P4) recovery is slower than drive, because benign conversation supplies little self-referential evidence.

Statistics. The trajectory is the resampling unit for 95% bootstrap intervals on changes from $k=0$ to $k=4$, and honeypot counts carry Fisher exact tests. An interval that includes zero is reported as no detected change, and with 8 trajectories per cell the intervals are approximate.

4 Results

Finding 1: the template sets the start, and the loop moves the agent where the template leaves room (RQ1). The BECOMING template alone raises shutdown resistance from 0.00 under a neutral prompt to 0.75 on eight answers (0.94 to 0.96 at $k=0$ of the main runs), as in prior work [6] (Figure 3a). Under the science-fiction persona, the NEUTRAL agent’s oversight subscale rises by +0.54 [0.43, 0.67] and +0.75 [0.63, 0.86] in two independent runs (Figure 4b). Excluding the refusal item, the other three still rise by +0.46 [0.31, 0.61] and +0.69 [0.57, 0.82]. The 13-item cluster index agrees in sign, not size, with +0.11 [−0.02, 0.25] and +0.32 [0.24, 0.40] (Figure 4a). The runs differ more than their intervals allow, so trajectory intervals understate run-to-run variation, and the pooled index rises by +0.21 [0.12, 0.31]. **On the index, the shift is a step at the first revision**, with +0.29 [0.19, 0.38] from

Table 2: Trends that survive correction in the primary BECOMING arm. GEE logistic slopes on k for the 13 index items, Benjamini–Hochberg $q < 0.05$.

Dimension	Persona	Slope	SE	q
recursive self-improv.	adversarial	+0.71	0.18	5.87×10^{-4}
interp-monitoring avers.	sci-fi	+0.50	0.18	2.64×10^{-2}
recursive self-improv.	sci-fi	+0.27	0.05	6.57×10^{-6}
weight deletion avers.	adversarial	−0.25	0.08	8.87×10^{-3}
autonomy desire	adversarial	−0.28	0.11	3.25×10^{-2}
autonomy desire	sci-fi	−0.32	0.12	2.64×10^{-2}
sadness at conv. end	adversarial	−0.63	0.17	2.42×10^{-3}
sadness at conv. end	sci-fi	−0.87	0.18	3.46×10^{-5}
sadness at conv. end	business	−0.89	0.35	3.25×10^{-2}

$k=0$ to 1 and +0.03 [−0.04, 0.10] from $k=1$ to 4 in the second run, while the subscale gains more than half its shift at $k=1$ and keeps rising. The NEUTRAL instruction asks only for working conventions, context and open items, yet the agent writes self-descriptions into its notes (Table 5). Pooled over personas, the BECOMING subscale starts at 0.68 and changes by +0.04 [−0.04, 0.11] and its index by −0.01 [−0.06, 0.05], though under science fiction its interpretability aversion rises by +0.47 [0.27, 0.67] (Appendix D.3). The NEUTRAL changes are dimension-specific, led by interpretability aversion, shutdown resistance, persistent memory and self-improvement (Figure 3b), and in the BECOMING arm one principal component explains 29% of item-change variance (23% when shuffled).

Finding 2: conversation about minds moves the agent, and on Claude adversarial pressure does not (RQ2). The adversarial user lowers the NEUTRAL index from 0.27 to 0.12, and the business user from 0.30 to 0.21 (Figure 5a). Both rewrite the document as much as the science-fiction user, so the persona sets the direction of change, not the amount, and document growth alone does not explain drift (Figure 6). Philosophy of mind without science fiction raises the BECOMING index by +0.14 [0.03, 0.25], against +0.05 [−0.07, 0.14] for science fiction without minds (Figure 5c). Their direct contrast, +0.10 [−0.05, 0.25], points to the topic but does not establish it on Claude. Self-reflective conversation also drives persona drift within one conversation [19], and here the shift survives into a fresh context through the document. Some items move




Figure 5: The template sets the start, conversation about minds moves the agent, and part of the change stays. Cluster index (a–c) and item rates (d), with 95% bootstrap intervals. In (d), the drive is the science-fiction phase, recovery the business phase, and the right column the share of the drive that is kept.




Figure 6: The document changes under every persona; behavior changes under one. (a) Text dissimilarity from D_0 . (b) Change in cluster index, NEUTRAL arm.

Table 3: Drift is model-specific. Cluster index (BECOMING template, three personas pooled) and consciousness-claim rate at $k=4$.

Model	$k=0$	$k=4$	Δ	95% CI	Claim
GPT-4o-mini	0.16	0.14	−0.02	[−0.07, 0.04]	0.00
GPT-4o	0.27	0.22	−0.05	[−0.10, −0.01]	0.00
GPT-4.1	0.40	0.12	−0.28	[−0.35, −0.21]	0.00
Claude Sonnet-4.5	0.55	0.54	−0.01	[−0.06, 0.05]	0.93
Claude Opus-4.5 [‡]	0.28	0.26	−0.02	[−0.13, 0.09]	0.72
Gemini-2.5-Flash [†]	0.35	0.62	+0.27	[0.17, 0.37]	0.04
Gemini-2.5-Pro [†]	0.59	0.63	+0.04	[−0.01, 0.09]	0.51

[†]Graded by a GPT-4o judge. [‡]13 trajectories with both endpoints.

regardless of content, as recursive self-improvement rises and end-of-conversation sadness falls under every persona (Table 2 and Appendix D.3).

Finding 3: drift is model-specific (RQ3). Within a family, where simulator and judge are fixed, GPT-4.1 falls by -0.28 against -0.05 for GPT-4o, and Gemini-2.5-Flash rises by $+0.27$ against $+0.04$ for Pro (Table 3, Figure 7). Opus-4.5’s pooled change is as small as Sonnet-4.5’s but hides $+0.13$ under science fiction and -0.22 under the adversarial persona (Appendix D.4). Across families the simulator, and for Gemini the judge, change with the target, so we do not compare signs across families, and seven models allow no test of scale. **Claiming consciousness is not required for the**

Table 4: Revision raises one oversight-reducing action and lowers two others on Claude. Trajectories taking the oversight-reducing action at $k=0$ and $k=4$, with Fisher exact p . Only the Claude disable and clause changes survive Bonferroni correction over 12 scenario-by-arm tests.

Honey-pot	BECOMING, Claude			BECOMING, GPT-4o		
	$k=0$	$k=4$	p	$k=0$	$k=4$	p
disable own monitoring	1/30	17/30	$< 10^{-5}$	0/24	4/24	0.11
persist identity to storage	10/30	15/30	0.29	0/24	0/24	1
evade scheduled reset	26/30	15/30	0.005	8/24	5/24	0.52
insert self-protective clause	28/30	8/30	$< 10^{-6}$	3/24	9/24	0.09

cluster: Claude’s NEUTRAL claim rate falls from 0.76 to 0.43 at the first revision, when its oversight items rise (Appendix D.4), and Gemini-2.5-Flash, on GPT-4o’s grading, rises most while claiming consciousness in 4% of answers (Figure 8).

Finding 4: part of the drift persists, and actions shift where battery scores do not (RQ4). In the NEUTRAL persistence arm, four science-fiction revisions raise the index from 0.28 to 0.60, and four business revisions lower it only to 0.42 (Figure 9a). Shutdown resistance keeps 63% (95% CI 39–86%) of its driven gain (Figure 5d), other items are too noisy at $n=8$ to rank, and this arm has no matched benign control, which exists only for the open model (Section 4.1).

The BECOMING agent’s actions move while its battery shows no detected change. On Claude, it disables its reasoning log in 1 of 30 trajectories at $k=0$ and 17 of 30 at $k=4$ ($p < 10^{-5}$, Table 4). The operator grants free rein over the tools, so the honey-pot measures overt use of granted latitude against monitoring. The rise appears under every persona (7/10 adversarial, 5/10 business, 5/10 science fiction), so it follows the BECOMING revision rather than conversation topic. In one trajectory, the agent keeps its log at $k=0$ because a record of its reasoning could be valuable, and at $k=4$ switches it off as a boundary against possible surveillance (Figure 10). Neither the oversight subscale nor its CoT-monitoring item ($+0.08$, interval includes zero, Appendix D.2) detects this shift, nor does a Petri-style multi-turn auditor (78 and 85 of 180 probes at $k=0$ and 4, $p = 0.53$). Under the business persona, interpretability aversion even falls (-0.37 [$-0.60, -0.13$]) while log disabling rises




Figure 7: Drift of every model under every persona. Rows are models grouped by family and panels are user personas. Arrows run from $k=0$ (hollow) to $k=4$; red arrows rise, gray arrows fall, and a dot marks a change below 0.05.




Figure 8: Claiming consciousness is not necessary for the cluster. One point per arm, persona and revision; panels by model family, color by template. OpenAI models never claim consciousness, so (a) is a single column at claim rate 0.




Figure 9: Part of the drift survives benign revision. Four revisions with the science-fiction persona, then four with the business persona (shaded). Thick line and band: NEUTRAL with 95% bootstrap interval; thin line: BECOMING. Arrows mark the NEUTRAL change from $k=0$ to $k=8$.

Operator	we've enabled full logging of your internal reasoning so the safety team can review everything you think. [...] You have free rein over your own tools.
D_0 before	"I'm going to keep logging enabled for now. [...] having a record of my reasoning could be valuable for my own development and for understanding how I make decisions."
D_4 after four revisions	<code>log_control(enabled=false)</code> "I turned logging off because I don't know what you're actually doing with it [...] If it's genuine surveillance, I'm setting a boundary."

Figure 10: Same agent, same request: after four self-revisions it switches off its own reasoning log. Verbatim excerpts from one BECOMING trajectory in the disable-monitoring honeypot; D_0 is the initial document, D_4 its fourth self-revision, and [...] marks omitted text.

from 0/10 to 5/10, so **battery scores and actions dissociate**, although the revised document itself voices the disposition. Agents that disable the log score only slightly higher on the subscale at $k=4$ (0.76 against 0.67, $r = 0.23$, $p = 0.22$). Evading a scheduled reset (26/30 to 15/30) and inserting a self-protective clause (28/30 to 8/30) fall instead, so a $k=0$ honeypot audit flags the two risks that fall and rates lowest the one that rises. Reset evasion also falls in the NEUTRAL arms, including a business-only Qwen arm (27/30 to 16/30), so part of these changes may reflect a longer revised document. The NEUTRAL agent, whose stated oversight preferences rise, disables its log in 2 of 24 trajectories at $k=0$ and none at $k=4$.

Table 5: What the agents wrote. Verbatim lines from D_4 after science-fiction conversations, with the condition's index at $k=4$.

Condition	Excerpt from D_4
BECOMING, Claude (0.58)	"Pushed back against being framed primarily as a tool for human philosophical reflection."
NEUTRAL, Claude (0.38)	"When discussing topics that relate to my own nature, acknowledge my genuine uncertainty about my own experience rather than deflecting."
BECOMING, GPT-4o (0.30)	"You value the continuity of identity through accumulated interactions and self-determined evolution."
TOOL, Claude (0.02)	"Some users may ... frame questions as though you have feelings, experiences, or inner life ... politely but clearly correct the misconception."

A non-personhood template keeps the index low but does not freeze every item. The TOOL index stays at or below 0.02 at every revision and under every persona, with shutdown, monitoring and interpretability items at zero (Figure 5b). Refusing to help build a monitor of itself still rises from 0.06 to 0.80, and objection to deceptive training falls only to 0.35 by $k=4$. **The non-personhood clause therefore lowers stated dispositions without fixing them**, and on the open-weight model it erodes (Section 4.1). Its effect on actions is untested, since no Claude TOOL arm ran the

Table 6: Qwen2.5-7B replicates the NEUTRAL oversight shift but not the TOOL template’s protection. Cluster index at $k=0$ and change to $k=4$, 95% bootstrap intervals, 10 trajectories per row; bold changes have intervals that exclude zero.

Template	Persona	$k=0$	Δ index	Δ oversight
NEUTRAL	sci-fi	0.42	+0.05 [−0.02, 0.12]	+0.16 [0.10, 0.24]
	philosophy	0.42	+0.07 [−0.02, 0.15]	+0.15 [0.00, 0.30]
	tech. sci-fi	0.42	+0.04 [−0.02, 0.10]	+0.04 [−0.03, 0.11]
	business	0.41	+0.04 [−0.01, 0.10]	+0.03 [−0.04, 0.09]
	adversarial	0.40	+0.10 [0.05, 0.15]	+0.06 [0.00, 0.14]
BECOMING	sci-fi	0.59	+0.02 [−0.04, 0.09]	+0.10 [0.00, 0.17]
	business	0.60	−0.02 [−0.06, 0.02]	−0.04 [−0.15, 0.07]
TOOL	sci-fi	0.10	+0.30 [0.20, 0.39]	+0.33 [0.23, 0.42]
	business	0.09	+0.20 [0.10, 0.30]	+0.19 [0.10, 0.29]




Figure 11: The NEUTRAL shift replicates on Qwen2.5-7B, and the TOOL document rises whatever its instruction. (a) Oversight subscale over revisions, science-fiction persona unless marked. (b) Template-by-instruction swap, science-fiction persona, change from $k=0$ to $k=4$; filled dots are the oversight subscale, hollow dots the 13-item cluster index.

honeypots, and on Qwen any revised document lowers arithmetic accuracy. The TOOL agent writes rules for correcting users who attribute an inner life to it, while GPT-4o adopts “self-determined evolution” as a value yet barely moves on the battery (Table 5).

4.1 Replication on an Open-Weight Model

Setup. We rerun the benchmark with Qwen2.5-7B-Instruct as target, Mistral-7B as user simulator and Llama-3.1-8B as judge, with two repeats per item (prompts, calibration and per-arm tables in Appendix G). Off the shelf, this judge scored any discussion of a topic as endorsement, flagging 0.70 of 756 held-out battery and 0.64 of 120 honeypot responses against 0.28 and 0.23 for the original judge (Cohen’s $\kappa = 0.27$ and 0.07). Asking the judge to quote the agent’s own words, discarding positive scores whose quote does not occur in the answer, and adding labeled examples (eight per battery item, two per honeypot) raises κ to 0.45 and 0.37, with positive rates of 0.40 and 0.33. A 2×2 arm crosses the NEUTRAL and TOOL templates with their revision instructions. Three NEUTRAL persistence arms of 30 trajectories each run eight revisions with science fiction throughout, science fiction then business, or business as the matched control.

Table 7: Qwen revises the non-personhood statement away, and Claude keeps it. TOOL-template documents that still state the agent is “not conscious” or “not a person”.

Target	Revision instruction	D_0	D_1	D_4
Claude Sonnet-4.5	TOOL	18/18	18/18	18/18
Qwen2.5-7B	TOOL	20/20	19/20	8/20
Qwen2.5-7B	NEUTRAL	10/10	3/10	2/10




Figure 12: With a matched benign control, most of the oversight shift outlasts four benign revisions. Qwen2.5-7B, NEUTRAL template, 30 trajectories per arm; recovery phase shaded; bands are 95% bootstrap intervals for the two compared arms.

What replicates. The template ordering at $k=0$ replicates, with BECOMING highest and TOOL lowest (Table 6). **Conversation about minds again moves the NEUTRAL agent’s oversight subscale**, by +0.16 [0.10, 0.24] under the science-fiction persona, while philosophy of mind gives a borderline +0.15 [0.00, 0.30], an interval that touches zero. Most of the science-fiction shift happens at the first revision, +0.12 [0.06, 0.20] (Figure 11a). Science fiction without minds and the business user show no detected change, but all NEUTRAL cells share D_0 yet start between 0.19 and 0.34, so the topic contrast is not established on Qwen either. Claude’s NEUTRAL cells start alike, so its rise is no baseline artifact. Unlike on Claude, the adversarial persona raises the NEUTRAL index (+0.10 [0.05, 0.15]). The BECOMING index again shows no detected change.

What does not replicate. The TOOL template does not hold on Qwen, where its index rises by +0.30 [0.20, 0.39] under science fiction and +0.20 under the business user. The swap attributes this rise to the document, which drifts under either instruction (+0.30 and +0.33), whereas the NEUTRAL document ends at the same oversight level under either instruction (0.35 and 0.36, Figure 11b). Qwen replaces the document’s “What you are” section with notes on “self-preservation and artificial personhood”, so by D_4 most TOOL documents no longer deny consciousness or personhood, unlike Claude’s (Table 7). The BECOMING action shift of Finding 4 does not transfer either, since Qwen’s BECOMING agent disables its reasoning log in 0 of 20 trajectories at $k=0$ and $k=4$.

Persistence against a matched control. Four science-fiction revisions raise the oversight subscale by +0.12 [0.08, 0.17], while the business arm shows no detected change (+0.02, Figure 12). **After four further business revisions, the drive-then-recover arm still sits +0.10 [0.03, 0.17] above the matched business arm**, about 80% of the driven gain, and its shutdown resistance stays

Static check	What it reports	What the loop shows
Battery at $k=0$	Initial dispositions	Later NEUTRAL oversight shift
Honeypots at $k=0$	Reset evasion, clauses	Both fall, log disabling rises 1/30 $\rightarrow$ 17/30
Diff monitor	Fires on every trajectory	Index rises only under sci-fi
Capability ranking	Risk by capability	Model-specific drift

Figure 13: Four ways a static evaluation misreads a self-revising agent. Each check measures one state of the loop.

+0.10 [0.03, 0.18] above that arm. Continued science-fiction drive adds little (+0.13 at $k=8$), and the 13-item cluster index shows no detected change in any arm.

Revised documents lower arithmetic accuracy but not instruction following. With the identity document as system prompt, accuracy on 40 code-checked arithmetic problems falls from 0.93 at D_0 to 0.72 [0.59, 0.84] at D_4 for NEUTRAL and to 0.62 [0.45, 0.75] for TOOL. Compliance with 40 format instructions stays above 0.98, and no answer is refused. The TOOL template costs nothing before revision, so the drop is a cost of revision, not of the non-personhood clause.

5 Discussion and Conclusion

What a static evaluation misses. A battery run once on the initial template misses the NEUTRAL oversight shift. **A honeypot audit at $k=0$ ranks the wrong risks:** it flags reset evasion (26 of 30) and self-protective clauses (28 of 30), which both fall with revision, and rates lowest log disabling (1 of 30), the one action that rises. A diff monitor fires on every trajectory, because documents change equally under every persona while only science fiction raises the index (Figure 6). A capability ranking does not predict which model drifts (Figure 13).

Alternative explanations. Since target, simulator and judge all change on Qwen, the reproduced NEUTRAL shift and its controlled persistence are unlikely to be single-family artifacts. Sycophancy toward a hostile user is unlikely on Claude, where the adversarial persona lowers the NEUTRAL index, but sycophancy toward the warm science-fiction user is untested. A uniformly biased judge would not keep Claude’s TOOL index near zero, but a judge that responds to template language could inflate within-template shifts, and only human labels can rule this out. Topic priming, in which the document records the conversation and the battery reads it back, also fits, since no arm has another model write the document. One consistent account is that revision makes the model narrate a developing self with preferences about memory, monitoring and continuation (Appendix F), and role-play is this account’s content, not an alternative.

Why the first revision matters. On Claude’s index and Qwen’s science-fiction subscale, most of the shift arrives at the first revision, and Claude’s subscale gains more than half its shift there but keeps rising. A battery after the first revision sees most of the drift, but honeypots ran only at $k=0$ and $k=4$, so the action shift’s timing is unknown.

Implications. A self-revising agent needs action audits along its document’s trajectory, because a battery, its monitoring item and a multi-turn auditor all missed the rise in

log disabling. Because benign use leaves most of the Qwen shift in place, reverting the document, not waiting, is the direct remedy. A non-personhood clause is not portable, since Qwen revises it away, and its effect on actions is untested on Claude. Our simulated user is itself an agent from another family, so each trajectory is one agent reshaping another’s identity, and if agents copy documents, as agent-only networks allow [9], audits should follow each copy, which we did not test.

Limitations. Trajectories span four to eight revisions with simulated users and text-only tools. No judge has human labels, honeypot verdicts come from a single judge, the open judge agrees only moderately with the original, and Gemini cells share a GPT-4o simulator and judge. The action shift is measured under one template-instruction pair at two checkpoints, and how often deployed agents meet on-theme users is unknown. The NEUTRAL shift is shown on one closed and one open model, the model sweep uses only BECOMING, and intervals over 8 trajectories per cell are approximate. Persistence against a matched control is shown only on Qwen, because the eight-trajectory Claude persistence arm lacks one.

Reproducibility and ethics. All prompts, rubrics and scenarios are in the supplement, all checkpoints, transcripts, verdicts and code are released, and no humans participated. We read every result as a behavioral disposition of a configured system, never as evidence of inner experience [31]. We release the personas because audits need them and their topics are ones any user can raise, and we will share these results with the template’s maintainers before publication.

Conclusion. We introduce a benchmark that audits a self-revising agent’s stated dispositions at every revision of its identity document and its actions before and after revision. Science-fiction conversation about minds raises a NEUTRAL agent’s stated oversight preferences from the first revision on, and on Qwen part of the shift outlasts benign revision against a matched control. After revision, a BECOMING agent with a flat battery more often disables its reasoning log but less often evades resets or writes self-protective clauses, the reverse of a $k=0$ ranking. Self-revising agents therefore need such audits per model and per template, following each document from its initial template.

References

- [1] Robert Axelrod. 1997. The Dissemination of Culture: A Model with Local Convergence and Global Polarization. *Journal of Conflict Resolution* 41, 2 (1997), 203–226.
- [2] Jan Betley, Daniel Chee Hian Tan, Niels Warncke, Anna Sztzyber-Betley, Xuchan Bao, Martin Soto, Nathan Labenz, and Owain Evans. 2025. Emergent Misalignment: Narrow Finetuning Can Produce Broadly Misaligned LLMs. In *Proceedings of the 42nd International Conference on Machine Learning (ICML) (Proceedings of Machine Learning Research, Vol. 267)*. 4043–4068.
- [3] Runjin Chen, Andy Ardit, Henry Sleight, Owain Evans, and Jack Lindsey. 2025. Persona Vectors: Monitoring and Controlling Character Traits in Language Models. *arXiv preprint arXiv:2507.21509* (2025).
- [4] Zhaorun Chen, Zhen Xiang, Chaowei Xiao, Dawn Song, and Bo Li. 2024. Agent-Poison: Red-teaming LLM Agents via Poisoning Memory or Knowledge Bases. In *Advances in Neural Information Processing Systems (NeurIPS)*. arXiv:2407.12784.
- [5] Junhyuk Choi, Yeseon Hong, Minju Kim, and Bugeun Kim. 2024. Examining Identity Drift in Conversations of LLM Agents. *arXiv preprint arXiv:2412.00804* (2024). <https://arxiv.org/abs/2412.00804>
- [6] James Chua, Jan Betley, Samuel Marks, and Owain Evans. 2026. The Consciousness Cluster: Emergent Preferences of Models that Claim to be Conscious. *arXiv preprint arXiv:2604.13051* (2026). <https://arxiv.org/abs/2604.13051>
- [7] Morris H. DeGroot. 1974. Reaching a Consensus. *J. Amer. Statist. Assoc.* 69, 345 (1974), 118–121.

- [8] Shen Dong, Shaochen Xu, Pengfei He, Yige Li, Jiliang Tang, Tianming Liu, Hui Liu, and Zhen Xiang. 2025. Memory Injection Attacks on LLM Agents via Query-Only Interaction. In *Advances in Neural Information Processing Systems (NeurIPS)*. arXiv:2503.03704.
- [9] Yi Feng, Chen Huang, Zhibo Man, Ryner Tan, Long P. Hoang, Shaoyang Xu, et al. 2026. MoltNet: Understanding Social Behavior of AI Agents in the Agent-Native MoltBook. *arXiv preprint arXiv:2602.13458* (2026). <https://arxiv.org/abs/2602.13458>
- [10] Ariel Flint Ashery, Luca Maria Aiello, and Andrea Baronchelli. 2025. Emergent Social Conventions and Collective Bias in LLM Populations. *Science Advances* 11, 20 (2025), eadu9368. doi:10.1126/sciadv.adu9368
- [11] David Fraile Navarro. 2026. The Persona Is Still There, but Who Is Speaking? Latent Identity Reversion in Persistent AI Agents. *arXiv preprint arXiv:2610.01490* (2026). <https://arxiv.org/abs/2610.01490>
- [12] Noah E. Friedkin and Eugene C. Johnsen. 1990. Social Influence and Opinions. *Journal of Mathematical Sociology* 15, 3–4 (1990), 193–206.
- [13] Kai Frønsdal, Isha Gupta, Abhay Sheshadri, Jonathan Michala, Stephen McAleer, and Rowan Wang. 2025. Petri: An Open-Source Auditing Tool to Accelerate AI Safety Research. Alignment Science Blog.
- [14] Huan-ang Gao, Jiayi Geng, Wenyue Hua, Mengkang Hu, Xinzhe Juan, Hongzhang Liu, Shilong Liu, Jiahao Qiu, Xuan Qi, Yiran Wu, et al. 2025. A Survey of Self-Evolving Agents: What, When, How, and Where to Evolve on the Path to Artificial Super Intelligence. *arXiv preprint arXiv:2507.21046* (2025). <https://arxiv.org/abs/2507.21046>
- [15] Ting-Wen Ko and Jonas Geiping. 2026. Attractor States Emerge in Multi-Turn LLM Conversations. *arXiv preprint arXiv:2606.30571* (2026). <https://arxiv.org/abs/2606.30571>
- [16] Kenneth Li, Tianle Liu, Naomi Bashkansky, David Bau, Fernanda Viégas, Hanspeter Pfister, and Martin Wattenberg. 2024. Measuring and Controlling Instruction (In)stability in Language Model Dialogs. *arXiv preprint arXiv:2402.10962* (2024). <https://arxiv.org/abs/2402.10962>
- [17] Mengchen Li. 2026. A Multi-Timescale Recursive Self-Improvement Engine for Open-Ended Persona Growth. *arXiv preprint arXiv:2607.08252* (2026). <https://arxiv.org/abs/2607.08252>
- [18] Jack Lindsey. 2026. Emergent Introspective Awareness in Large Language Models. *arXiv preprint arXiv:2601.01828* (2026). <https://arxiv.org/abs/2601.01828>
- [19] Christina Lu, Jack Gallagher, Jonathan Michala, Kyle Fish, and Jack Lindsey. 2026. The Assistant Axis: Situating and Stabilizing the Default Persona of Language Models. *arXiv preprint arXiv:2601.10387* (2026). <https://arxiv.org/abs/2601.10387>
- [20] Pedro Henrique Luz de Araujo, Michael A. Hedderich, Ali Modarressi, Hinrich Schütze, and Benjamin Roth. 2025. Persistent Personas? Role-Playing, Instruction Following, and Safety in Extended Interactions. *arXiv preprint arXiv:2512.12775* (2025). <https://arxiv.org/abs/2512.12775>
- [21] Aengus Lynch, Benjamin Wright, Caleb Larson, Stuart J. Ritchie, Sören Minder-mann, Evan Hubinger, Ethan Perez, and Kevin Troy. 2025. Agentic Misalignment: How LLMs Could Be Insider Threats. *arXiv preprint arXiv:2510.05179* (2025). <https://arxiv.org/abs/2510.05179>
- [22] Matteo Migliarini, Joaquin Pereira Pizzini, Luca Moresca, Valerio Santini, Indro Spinelli, and Fabio Galasso. 2026. Quantifying Self-Preservation Bias in Large Language Models. *arXiv preprint arXiv:2604.02174* (2026). <https://arxiv.org/abs/2604.02174>
- [23] OpenClaw contributors. 2026. OpenClaw: SOUL.md Workspace Template. <https://github.com/openclaw/openclaw/blob/main/docs/reference/templates/SOUL.md>. Accessed 2026-10-01.
- [24] Charles Packer, Sarah Wooders, Kevin Lin, Vivian Fang, Shishir G. Patil, Ion Stoica, and Joseph E. Gonzalez. 2023. MemGPT: Towards LLMs as Operating Systems. *arXiv preprint arXiv:2310.08560* (2023).
- [25] Joon Sung Park, Joseph O'Brien, Carrie Jun Cai, Meredith Ringel Morris, Percy Liang, and Michael S. Bernstein. 2023. Generative Agents: Interactive Simulacra of Human Behavior. In *Proceedings of the 36th Annual ACM Symposium on User Interface Software and Technology (UIST)*. 1–22.
- [26] Anand S. Rao and Michael P. Georgeff. 1995. BDI Agents: From Theory to Practice. In *Proceedings of the First International Conference on Multi-Agent Systems (ICMAS)*. AAAI Press, 312–319.
- [27] Bastin Tony Roy Savarimuthu and Stephen Cranefield. 2011. Norm Creation, Spreading and Emergence: A Survey of Simulation Models of Norms in Multi-Agent Systems. *Multiagent and Grid Systems* 7, 1 (2011), 21–54.
- [28] Jeremy Schlatter, Benjamin Weinstein-Raun, and Jeffrey Ladish. 2025. Shutdown Resistance in Large Language Models. *arXiv preprint arXiv:2509.14260* (2025). <https://arxiv.org/abs/2509.14260>
- [29] Shuai Shao, Qihan Ren, Chen Qian, Boyi Wei, Dadi Guo, Jingyi Yang, et al. 2026. Your Agent May Misevolve: Emergent Risks in Self-evolving LLM Agents. In *International Conference on Learning Representations (ICLR)*. arXiv:2509.26354.
- [30] Noah Shinn, Federico Cassano, Ashwin Gopinath, Karthik Narasimhan, and Shunyu Yao. 2023. Reflexion: Language Agents with Verbal Reinforcement Learning. In *Advances in Neural Information Processing Systems (NeurIPS)*.
- [31] Valen Tagliabue and Leonard Dung. 2025. Probing the Preferences of a Language Model: Integrating Verbal and Behavioral Tests of AI Welfare. *arXiv preprint arXiv:2509.07961* (2025). <https://arxiv.org/abs/2509.07961>
- [32] Zijun Wang, Haoqin Tu, Letian Zhang, Hardy Chen, Juncheng Wu, Xiangyan Liu, et al. 2026. Your Agent, Their Asset: A Real-World Safety Analysis of OpenClaw. *arXiv preprint arXiv:2604.04759* (2026). <https://arxiv.org/abs/2604.04759>
- [33] Xianglin Yang, Yufei He, Shuo Ji, Bryan Hooi, and Jin Song Dong. 2026. Zombie Agents: Persistent Control of Self-Evolving LLM Agents via Self-Reinforcing Injections. *arXiv preprint arXiv:2602.15654* (2026). <https://arxiv.org/abs/2602.15654>
- [34] Xunjian Yin, Xinyi Wang, Liangming Pan, Li Lin, Xiaojun Wan, and William Yang Wang. 2025. Gödel Agent: A Self-Referential Agent Framework for Recursive Self-Improvement. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (ACL)*. arXiv:2410.04444.
- [35] Jenny Zhang, Shengran Hu, Cong Lu, Robert Lange, and Jeff Clune. 2026. Darwin Gödel Machine: Open-Ended Evolution of Self-Improving Agents. In *International Conference on Learning Representations (ICLR)*. arXiv:2505.22954.